

TOWSON UNIVERSITY

MASTER'S PROJECT

Aiding Linear Television Media Planning Through Bayesian Inference and Forecasting

Author:
Matthew TIGER

Supervisors:
Mr. Jason MUHLENKAMP
Dr. M.D. VOISEI

*A project submitted in fulfillment of the requirements
for the degree of Master of Science*

at

Towson University

April 29, 2018

Declaration of Authorship

I, Matthew TIGER, declare that this project titled, “Aiding Linear Television Media Planning Through Bayesian Inference and Forecasting” and the work presented in it are my own. I confirm that:

- This work was done wholly or mainly while in candidature for a research degree at this University.
- Where any part of this project has previously been submitted for a degree or any other qualification at this University or any other institution, this has been clearly stated.
- Where I have consulted the published work of others, this is always clearly attributed.
- Where I have quoted from the work of others, the source is always given. With the exception of such quotations, this project is entirely my own work.
- I have acknowledged all main sources of help.
- Where the thesis is based on work done by myself jointly with others, I have made clear exactly what was done by others and what I have contributed myself.

Signed:

Date:

TOWSON UNIVERSITY

Abstract

Department of Mathematics

Master of Science

Aiding Linear Television Media Planning Through Bayesian Inference and Forecasting

by Matthew TIGER

The Thesis Abstract is written here (and usually kept to just this page). The page is kept centered vertically so can expand into the blank space above the title too...

Acknowledgements

The acknowledgments and the people to thank go here, don't forget to include your project advisor...

Contents

Declaration of Authorship	iii
Abstract	v
Acknowledgements	vii
1 Introduction	1
1.1 Background	1
1.2 Motivation	1
1.3 Challenges	2
1.4 Problem Description	2
2 Data	3
2.1 TV supplier data	3
2.1.1 Linear Media Schedule	3
2.1.2 Forecasted Demographic Estimates	3
2.2 Audience measurement data	4
2.3 Data Used	5
3 Model	7
3.1 Preliminaries	7
3.1.1 Units of analysis and observation	7
3.1.2 Covariates	7
3.2 Assumptions	8
3.3 Description	8
3.4 Prior Distribution Choice	8
3.5 Inference	8
3.6 Validation	9
3.7 Model selection and Sensitivity Analysis	9
4 Forecast Results	11
4.1 Predictive Accuracy	11
4.2 Media plans	11
4.2.1 Random Media Plans	11
4.2.2 Calibration	11
4.3 New Content	11
4.4 Inference Summary	11
5 Conclusion	13

List of Figures

List of Tables

2.1	Example linear media schedule data for two selling title airings.	3
2.2	Example forecasted demographic estimates data for two selling title weeks.	4
2.3	Example program table data for two historical program airings.	4
2.4	Example viewing table data for one respondent's viewing on historical program airings from Table 2.3.	4
2.5	Example data set used for model for two selling title airings.	6

List of Abbreviations

LAH List Abbreviations Here
WSF What (it) Stands For

For/Dedicated to/To my...

1. Introduction

The world of advertising today consists of

We are trying to provide better linear tv media plans.

We will use the theory of Bayesian statistics to develop a model that will forecast future telecast airings' audience impression concentration distributions corrected for small sample sizes present in the measurement training data.

1.1 Background

TV sellers have airtime available for sale to be used to air advertisements. TV buyers purchase airtime from sellers in order to air their desired advertisements.

TV buyers want to buy ad airtime that will help them share their message with a target audience.

A media plan consists of an allocation of a number of equivalized units to a collection of selling title / weeks.

Audience measurement companies provide measurements television viewership. Some companies take these measurements using a statistical sample of the TV viewing population and tracking the viewership of that sample and extrapolating to the population at large.

TV sellers bundle together airtime into a unit called a selling title. This unit contains logically grouped airings of content. Sometimes the content within the selling title is specific, other times it's a blend of content packaged together.

An example of a piece of content available is *Funny Animated Cartoon* and an example of a selling title would be *Funny Animated Cartoon M-F 8PM*.

TV sellers accept media plans allocated at the selling title/week level and then as the flight date approaches determine the actual ad schedule that will air. This is known as trafficking the ads.

The planning is done based off of demo estimates, which are impression estimates of a particular age/gender demographic. The TV seller provided these estimates using their substantive knowledge. The planning for strategic targets is done by forecasting a percentage. The resulting forecast is then the percentage multiplied by the forecasted demo estimate

1.2 Motivation

The sample sizes used for the majority of telecasts in the measurement source are too small to be trusted on their own. This is the major challenge of providing forecasts for media plans; using the data on its own will produce inaccurate forecasts that do not take into account the small sample size.

(Show picture of sample sizes in networks)

A motivating example is related to evaluating the performance of baseball players: players are judged by their batting average but this metric is not informative with few at-bats. With more information about the league and past historical performances we can use bayesian inference to come up with a better estimate that takes such factors into account as well as the observed at-bats.

Simiarly, we hope to use the data from past airings to perform better inference on telecasts with small sample sizes.

1.3 Challenges

Low sample size must be addressed. Variance within selling-title/week. Impression estimates are tied to demo forecasts.

For the purposes of this thesis we will be addressing the first issue alone. We will assume for testing that we know the outcome of the demo with prescience and that we can traffick the ads wherever we want.

1.4 Problem Description

Summarize the challenges and what we are trying to solve. Specifically, given the “noisy” measurement data and the supplier data, can we create model to forecast the in-target impressions for future selling title airings with some degree of accuracy.

2. Data

In order to provide the forecasts necessary for media planning, we make use of several data sets. Namely, TV supplier data, which contains the inventory to plan against, and audience measurement data, which contains historical program airings and the amount of viewing associated to those airings for various audiences. We discuss these data sets in more detail in the sections below.

2.1 TV supplier data

TV suppliers provide us with a holistic view of their inventory that is for sale. The most granular unit in their inventory is a selling title during a future broadcast week, i. e. a week that starts on a Monday and ends the following Sunday. This unit is a grouping of content that will air during a given broadcast week and the grouping can be based on either similar time of airing or similar content. The two most important pieces of data provided by TV suppliers are the linear media schedule, which contains the exact air times of selling titles, and the forecasted demographic estimates, both of which are explained below.

2.1.1 Linear Media Schedule

The linear media schedule lists all scheduled program airings on a given network. Associated to each airing is an identifier for the selling title and an identifier for the content. Note that the selling title indicates the grouping of content that is sold while the content indicates the nature of the programming that will be shown. Example records from a linear media schedule are shown in Table 2.1.

network	selling title	selling title name	content	content name	start datetime	end datetime
BCST	100	Adult Cartoon 8PM	10	Adult Cartoon	2017-04-02 20:00:00	2017-04-02 20:30:00
BCST	101	Adult Cartoon 8:30PM	10	Adult Cartoon	2017-04-02 20:30:00	2017-04-02 21:00:00

TABLE 2.1: Example linear media schedule data for two selling title airings.

As can be seen from Table 2.1, the same content identifier can be associated to multiple selling titles. This indicates that the TV supplier wishes to differentiate the content airing at 8:00 PM from the content airing at 8:30 PM, usually to price the airings differently.

2.1.2 Forecasted Demographic Estimates

The forecasted demographic estimates data contains for each selling title, broadcast week, and demographic audience the expected impressions per unit, i. e. the expected number of audience members that will watch a commercial airing during the given selling title and broadcast week.

Example records from the forecasted demographic estimates data are shown in Table 2.2

selling title	broadcast week	demographic	impressions per unit
100	2017-03-27 06:00:00	F45-49	150000
100	2017-03-27 06:00:00	P18-49	1500000
101	2017-03-27 06:00:00	F45-49	120000
101	2017-03-27 06:00:00	P18-49	1000000

TABLE 2.2: Example forecasted demographic estimates data for two selling title weeks.

2.2 Audience measurement data

As mentioned previously, the audience measurement data contains the programs that have aired historically and who watched them according to the audience measurement source. This data is stored in two separate tables, the program table and the viewing table.

The program table details the historical program airings and meta data about those airings. A program in the table corresponds to an identifier of a logical grouping of aired content per the measurement source while the telecast corresponds to an identifier of a single airing of a program. Some information included is the network the program aired on, the start and end datetime of the program airing, as well as the genre associated to the program and whether or not this is a repeat airing of the telecast. Note that the genre information serves to categorize the programs aired by grouping together programs with similar content. Some example program data is provided in Table 2.3.

network	program	telecast	program name	start datetime	end datetime	genre id	is first run	is live
BCST	1000	301	Adult Cartoon	2017-04-02 20:02:00	2017-04-02 20:30:00	Animation	1	0
BCST	1000	302	Adult Cartoon	2017-04-02 20:30:00	2017-04-02 21:00:00	Animation	1	0

TABLE 2.3: Example program table data for two historical program airings.

Associated with the program table is the viewing table which details the respondents' viewing of historical programs. This data is stored at the minute level and contains which respondents were watching each minute of a given telecast and how many commercial seconds aired during the minute as well as the total number of commercial seconds that aired. Additionally, limited identifying information about the respondent is stored in this table such as the respondent's age and gender. Some example viewing data is provided in Table 2.4.

program	telecast	respondent	minute	comm secs	weight	age	gender	total comm secs
1000	301	2	3	60	2050	48	F	120
1000	301	2	4	45	2050	48	F	120
1000	301	2	15	60	2050	48	F	120
1000	301	2	16	30	2050	48	F	120
1000	302	2	22	15	2050	48	F	100
1000	302	2	23	60	2050	48	F	100

TABLE 2.4: Example viewing table data for one respondent's viewing on historical program airings from Table 2.3.

Of primary interest from this data set is the impressions per unit, or average commercial minute (ACM), for a given audience. The definition of a program airing's ACM measurement is the weighted average of the in-target commercial viewing seconds over the number of commercial seconds that aired during the program's airing.

Formally, let A be the set of panelists that are considered in-target out of a total of n possible panelists. Suppose that program airing i has t_i minutes of programming and let s_{ij} be the number of commercial seconds during minute j of the airing. Further, let $p_{ijk} = 1$ if panelist k was watching during the j -th minute of program airing i and 0 otherwise and let w_k be the weight assigned to the panelist by the measurement source. Then, m_i^A , the ACM of program airing i is

$$m_i^A = \left\lceil \frac{\sum_{k=1}^n w_k \sum_{j=1}^{t_i} s_{ij} p_{ijk} \mathbf{1}_A(p_k)}{\sum_{j=1}^{t_i} s_{ij}} \right\rceil \quad (2.1)$$

A related quantity for program airing i , the impression concentration of target A relative to target B where $A \subseteq B$, is

$$c_i = \frac{m_i^A}{m_i^B}.$$

This quantity is used to measure the proportion of impressions attributed to target A relative to target B .

As an example, if we consider target A to be the set containing only respondent 2, then from Table 2.4 we see that for program 1000 and telecast 302, $m_{(1000,302)}^A = 1538$.

2.3 Data Used

For the purposes of this project, we limit the data used from the above data sets to three networks labeled BCST, ETMT, and SPTS. These networks are a broadcast network that has large reach, a cable entertainment network, and a cable sports network, respectively. Further, we limit the audience measurement and linear media schedule data to only airings that occur between 2015-12-28 06:00:00 and 2018-01-01 06:00:00 which coincides with broadcast years 2016 through 2017. We shall consider the in-target audience, target A , to be Females aged 45 through 49 and the demographic audience, target B , to be Persons aged 18 through 49. The demographic audience was chosen as it is representative of the demographic audience typically chosen when constructing a media plan while the in-target audience was chosen as it is small relative to the total number of TV viewers in the United States, roughly 3.3%. As a consequence, the size of the in-target panel in the audience measurement data is small and the target is exemplary of the issues this project aims to solve.

The data used for training and testing the model is combined from the above tables to create a data set that associates to each selling title airing in the linear media schedule the program airing in the measurement data whose start and end datetimes most overlap with the start and end datetimes in the linear media schedule. We then assign the selling title airing the ACM of the matched program airing for target A and target B . Example data is shown in Table 2.5.

network	selling title	content	start datetime	end datetime	program	telecast	ACM A	ACM B
BCST	100	10	2017-04-02 20:00:00	2017-04-02 20:30:00	1000	301	110560	1203560
BCST	101	10	2017-04-02 20:30:00	2017-04-02 21:00:00	1000	302	210560	1501000

TABLE 2.5: Example data set used for model for two selling title airings.

3. Model

3.1 Preliminaries

In the following section we will provide preliminary information necessary to understand the model.

3.1.1 Units of analysis and observation

We are most interested in the content that airs.

the units of observation are selling title airings in the schedule. The outcome variable we are measuring for each unit i is r_i^A the average number of panelists associated to target A given the average number of panelists associated to target B where $A \subseteq B$. Of primary interest is the proportion of panelists in target A relative to target B .

3.1.2 Covariates

Using the conclusions of paper TV ratings, we include the following covariates into the model:

1. Broadcast Month - The broadcast month associated to the start of a selling title airing. "The key link between the broadcast and Gregorian calendars is that the first week of every broadcast month always contains the Gregorian calendar first of the month."
2. Day of Week - The day of the week associated to the start of a selling title airing. These are encoded from 0 - 6 with 0 being Monday and 6 being Saturday
3. Stratified Hour - We define the following groupings of hours as this covariate. These groupings are adapted from the measurement source. morning - airing start hour is in (6, 7, 8, 9), daytime - airing start hour is in (10, 11, 12, 13, 14), early_fringe - airing start hour is in (15, 16, 17, 18), prime_1900_2000 - airing start hour is 19, prime_2000_2100 - airing start hour is 20, prime_2100_2200 - airing start hour is 21, prime_2200_2300 - airing start hour is 22, late_fringe - airing start hour is in (23, 0, 1), graveyard - airing start hour is in (2, 3, 4, 5)
4. Content Id - The id denoting the content associated with the scheduled airing. e.g. Famous Funny Animated Cartoon show
5. Lead-in Content Id - The id denoting the content associated with the preceding airing.
6. first-run - Denotes if the airing is not a repeat airing. This is a 1 if it is not a repeat and 0 if it is a repeat. This information is obtained through the measurement source.

7. genre-id - Denotes the genre associated to the airing. This conveys the general flavor of the aired content. We will enumerate the different genre-ids and their meaning. This information is obtained through the measurement source.

3.2 Assumptions

We assume that the observations are exchangeable given the parameters of the model and the covariates of the selling title airing. We describe the definition of exchangeable here. Exchangeability is important since it allows us to use Bayes' Theorem to compute the posterior. The number of draws from the binomial aren't independent, but we use this as an approximation. this assumption is imperfect.

3.3 Description

See below for a formal description of the model. Since the data we observe is the average number of panelists in target A given the average number of panelists in target B with $A \subseteq B$, it is natural to model the likelihood as a binomial distribution with unknown probability parameter π .

$$\begin{aligned}\theta_i &= \beta_0 + \sum_{j=1}^m X_i^\top \beta_j, \quad \beta_j \sim \mathcal{N}(0, 4^2) \\ \omega_i &= \text{logit}^{-1}(\theta_i) \\ \kappa_i &\sim \text{Exp}(X_i^\top \lambda), \quad \lambda \sim \mathcal{U}(1e-3, 1) \\ \pi_i &\sim \text{Beta}(\omega_i \kappa_i + 1, (1 - \omega_i) \kappa_i + 1) \\ y_i | n_i, \theta_i, X_i &\sim \text{Bin}(n_i, \pi_i)\end{aligned}$$

$$\text{where } \text{logit}^{-1}(\alpha) = \frac{\exp \alpha}{1 + \exp \alpha}.$$

3.4 Prior Distribution Choice

The coefficients are chosen so that the main mass of the distribution on the logit scale falls between $1e-5$ and $1 - 1e-5$.

We chose the lambda to be uniform on $1e-3$ and 1 since the kappa variable is the concentration parameter which is roughly the number of trials in the binomial likelihood. For any given telecast airing we will have at most the number of trials be the number of panelists which cannot exceed 100K. This prior on lambda and the exponential reflects that.

We chose to separate variance based off of is-first-run since we noticed that this variable is indicative of how many trials there will be for an airing.

3.5 Inference

We perform the inference using pymc3. We will display convergence checks such as gelman-rubin statistics, effective number of parameters among other things to ensure that posterior distribution obtained are representative of the actual posterior distribution.

Show the calculation of \hat{r} and \hat{n}_{eff} and display summary statistics of values to show that inferences have approximate convergence and that trace objects are representative of

3.6 Validation

posterior predictive checking.

Given the posterior distribution obtained through inference, we make draws from the distribution to arrive at replicated data, y^{rep} . We check the distribution of the replicated data using test statistics comparing replicated data to the observed test statistic in the data.

- We check the resulting posteriors using test statistics of the mean, min, max, and std of the number of in-target panelists. This makes sure that the data generated by the model are consistent with the observed data.
- ordered discrete data check to see if data is under/overdispersed compared to replicated data.
- binned residual plots against observed in-target panelists to check that the regression fits make sense.
- The data was collected sequentially in time so we must check if residuals have autocorrelation.
- outlier check to see if any impressions conc predictions are greater than 1-2 std of impressions conc.
- realized residual plots for impressions conc

3.7 Model selection and Sensitivity Analysis

We use the WAIC to determine the model that has the most predictive accuracy. The WAIC is chosen because it is Bayesian and able to be compared across models with same likelihood.

Go through the definition of WAIC and explain how it is computed etc.

Here we will also check to see if the choice in prior is affecting inference.

- Remove the hierarchical nature on the variance parameter and evaluate how inference is affected - change normal distribution on coefficients to cauchy or student t to see if inference is affected

4. Forecast Results

We want to compare the forecasts from the posterior predictive distribution on test data vs the industry standard forecasts per Tv ratings paper.

4.1 Predictive Accuracy

Compute probabilistic scores for bayes forecast and compute MAE for industry standard. Compare vs baseline of population proportion.

- Show actual vs predicted plots of each airing

- Walk through example that illustrates the added benefit of model

4.2 Media plans

4.2.1 Random Media Plans

Create media plans by randomly selecting airings and then compute HPD and mean of media plan where we assume that the demo estimates are known.

- We want to check that outcomes are within are 95% HPD

4.2.2 Calibration

Show calibration plot of binning up all airings into 20 or so bins and then plotting mean of actual vs mean of predicted. If model is calibrated then “low stuff should perform low and the high stuff high.”

4.3 New Content

We expect the model to perform better on new content than industry standard method. We should illustrate that here.

4.4 Inference Summary

We summarize the inferences using the mean of the posterior.

5. Conclusion